

## Spoken Digit Recognition

In this notebook, You will do Spoken Digit Recognition.

Input - speech signal, output - digit number

It contains

1. Reading the dataset. and Preprocess the data set. Detailed instructions are given below. You have to write the code in the same cell which contains the instruction.
2. Training the LSTM with RAW data
3. Converting to spectrogram and Training the LSTM network
4. Creating the augmented data and doing step 2 and 3 again.

### Instructions:

1. Don't change any Grader Functions. Don't manipulate any Grader functions. If you manipulate any, it will be considered as plagiarised.

2. Please read the instructions on the code cells and markdown cells. We will explain what to write.

Saving... 

Format what we asked. Eg. Don't return List of we are asking for a numpy array.

4. Please read the external links that we are given so that you will learn the concept behind the code that you are writing.

5. We are giving instructions at each section if necessary, please follow them.

### Every Grader function has to return True.

```
1 from google.colab import drive  
2 drive.mount('/content/drive')
```

Mounted at /content/drive

```
1 !pip install gdown  
2 !pip install librosa  
3 import numpy as np  
4 import pandas as pd  
5 import librosa  
6 import os  
7 import gdown  
8 import zipfile  
9 import os  
10 import sklearn  
11 from sklearn.model_selection import train_test_split  
12 import librosa  
13 import warnings  
14 warnings.filterwarnings("ignore")  
15 import matplotlib.pyplot as plt  
16 import tensorflow as tf  
17 from tqdm import tqdm  
18 from sklearn.utils import shuffle  
19 from tensorflow.keras.preprocessing.sequence import pad_sequences  
20 from tensorflow.keras.layers import Input, LSTM, Dense, Dropout  
21 from tensorflow.keras.models import Model  
22 from sklearn.metrics import roc_auc_score, f1_score  
23 from keras.callbacks import ModelCheckpoint, TensorBoard  
24 ##if you need any imports you can do that here.
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: gdown in /usr/local/lib/python3.8/dist-packages (4.4.0)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.8/dist-packages (from gdown) (4.6.3)
Requirement already satisfied: requests[socks] in /usr/local/lib/python3.8/dist-packages (from gdown) (2.25.1)
Requirement already satisfied: six in /usr/local/lib/python3.8/dist-packages (from gdown) (1.15.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.8/dist-packages (from gdown) (4.64.1)
Requirement already satisfied: filelock in /usr/local/lib/python3.8/dist-packages (from gdown) (3.9.0)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.8/dist-packages (from requests[socks]>gdown) (1.26.14)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.8/dist-packages (from requests[socks]>gdown) (2022.12.7)
Requirement already satisfied: chardet<5,>=3.0.2 in /usr/local/lib/python3.8/dist-packages (from requests[socks]>gdown) (4.0.0)
Requirement already satisfied: idna3,>=2.5 in /usr/local/lib/python3.8/dist-packages (from requests[socks]>gdown) (2.10)
Requirement already satisfied: PySocks!=1.5.7,>=1.5.6 in /usr/local/lib/python3.8/dist-packages (from requests[socks]>gdown) (1.7.1)
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: librosa in /usr/local/lib/python3.8/dist-packages (0.8.1)
Requirement already satisfied: numpy>=1.15.0 in /usr/local/lib/python3.8/dist-packages (from librosa) (1.22.4)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.8/dist-packages (from librosa) (23.0)
Requirement already satisfied: numba>=0.43.0 in /usr/local/lib/python3.8/dist-packages (from librosa) (0.56.4)
Requirement already satisfied: joblib>=0.14 in /usr/local/lib/python3.8/dist-packages (from librosa) (1.2.0)
Requirement already satisfied: resampy>=0.2.2 in /usr/local/lib/python3.8/dist-packages (from librosa) (0.4.2)
Requirement already satisfied: soundfile>=0.10.2 in /usr/local/lib/python3.8/dist-packages (from librosa) (0.12.1)
Requirement already satisfied: pooch>=1.0 in /usr/local/lib/python3.8/dist-packages (from librosa) (1.7.0)
Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.8/dist-packages (from librosa) (1.10.1)
Requirement already satisfied: decorator>=3.0.0 in /usr/local/lib/python3.8/dist-packages (from librosa) (4.4.2)
Requirement already satisfied: scikit-learn!=0.19.0,>=0.14.0 in /usr/local/lib/python3.8/dist-packages (from librosa) (1.2.1)
Requirement already satisfied: audioread>=2.0.0 in /usr/local/lib/python3.8/dist-packages (from librosa) (3.0.0)
Requirement already satisfied: importlib-metadata in /usr/local/lib/python3.8/dist-packages (from numba>=0.43.0->librosa) (6.0.0)
Requirement already satisfied: llvmlite<0.40,>=0.39.0dev6 in /usr/local/lib/python3.8/dist-packages (from numba>=0.43.0->librosa) (0.39.1)
Requirement already satisfied: setuptools in /usr/local/lib/python3.8/dist-packages (from numba>=0.43.0->librosa) (57.4.0)
Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.8/dist-packages (from pooch>=1.0->librosa) (2.25.1)
Requirement already satisfied: platformdirs>=2.5.0 in /usr/local/lib/python3.8/dist-packages (from pooch>=1.0->librosa) (3.0.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.8/dist-packages (from scikit-learn!=0.19.0,>=0.14.0->librosa) (3.1.0)
```

```
Requirement already satisfied: cffi>=1.0 in /usr/local/lib/python3.8/dist-packages (from soundfile>=0.10.2->librosa) (1.15.1)
Requirement already satisfied: pycparser in /usr/local/lib/python3.8/dist-packages (from cffi>=1.0->soundfile>=0.10.2->librosa) (2.21)
Requirement already satisfied: chardet<5,>=3.0.2 in /usr/local/lib/python3.8/dist-packages (from requests>=2.19.0->pooch>=1.0->librosa) (4.0.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.8/dist-packages (from requests>=2.19.0->pooch>=1.0->librosa) (2022.12.7)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.8/dist-packages (from requests>=2.19.0->pooch>=1.0->librosa) (2.10)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.8/dist-packages (from requests>=2.19.0->pooch>=1.0->librosa) (1.26.14)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.8/dist-packages (from importlib-metadata->numba>=0.43.0->librosa) (3.15.0)
```

We shared recordings.zip, please unzip those.

```
1 tf.test.gpu_device_name()
'/device:GPU:0'

1 !gdown 1BqrqsLqWvKa2m1Dw7KeWlME723VRpkTO

  Downloading...
  From: https://drive.google.com/uc?id=1BqrqsLqWvKa2m1Dw7KeWlME723VRpkTO
  To: /content/Copy of recordings.zip
  100% 9.28M/9.28M [00:00<00:00, 14.4MB/s]

1 zip_ref = zipfile.ZipFile("/content/drive/MyDrive/Copy of recordings.zip", "r")
2 zip_ref.extractall()
3 zip_ref.close()

Saving...  Saving...

1 #read the all file names in the recordings folder given by us
2 #(if you get entire path, it is very useful in future)
3 #save those files names as list in "all_files"
4 all_files = [file for file in tqdm(os.listdir("/content/recording"))]

100%|██████████| 2000/2000 [00:00<00:00, 840120.98it/s]
```

Grader function 1

```
1 def grader_files():
2     temp = len(all_files)==2000
3     temp1 = all([x[-3:]=="wav" for x in all_files])
4     temp = temp and temp1
5     return temp
6 grader_files()
```

True

Create a dataframe(name=df\_audio) with two columns(path, label).

You can get the label from the first letter of name.

Eg: 0\_jackson\_0 -> 0

0\_jackson\_43 -> 0

```
1 audio_data = [{'path':path,'label':path[0]} for path in tqdm(all_files)]
2 df_audio = pd.DataFrame(audio_data)

100%|██████████| 2000/2000 [00:00<00:00, 575626.71it/s]
```

```
1 df_audio.head()
```

	path	label
0	1_yweweler_6.wav	1
1	8_yweweler_48.wav	8
2	6_yweweler_35.wav	6
3	7_theo_16.wav	7
4	7_nicolas_40.wav	7

## Exploring the sound dataset

```
1 #It is a good programming practise to explore the dataset that you are dealing with. This dataset is unique in itself because it has sounds as input
2 #https://colab.research.google.com/github/Tyler-Hilbert/AudioProcessingInPythonWorkshop/blob/master/AudioProcessingInPython.ipynb
3 #visualize the data and write code to play 2-3 sound samples in the notebook for better understanding.
4 #please go through the following reference video https://www.youtube.com/watch?v=37zCgCdV468
```

```
1 # download the thinkdsp library
2 !git clone https://github.com/AllenDowney/ThinkDSP.git

Cloning into 'ThinkDSP'...
remote: Enumerating objects: 2469, done.
remote: Total 2469 (delta 0), reused 0 (delta 0), pack-reused 2469
Receiving objects: 100% (2469/2469), 208.82 MiB | 20.26 MiB/s, done.
Resolving deltas: 100% (1353/1353), done.
```

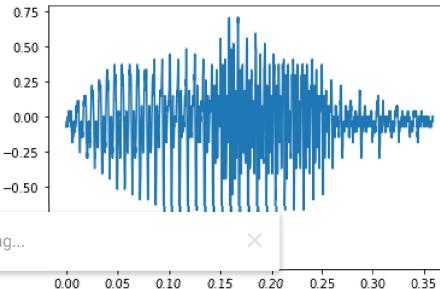
```
1 import sys
2 sys.path.insert(0, 'ThinkDSP/code/')
3 import thinkdsp
```

```
4 import matplotlib.pyplot as pyplot
5 import IPython
1 wave = thinkdsp.read_wave('/content/recording/0_nicolas_32.wav')
2 wave.play()
3 IPython.display.Audio('sound.wav')
```

Writing sound.wav

0:00 / 0:00

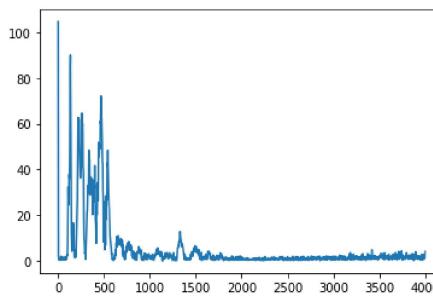
```
1 segment = wave.segment()
2 segment.plot()
3 pyplot.show()
```



Saving...

X

```
1 # Plot spectrum of wave
2 spectrum = wave.make_spectrum()
3 spectrum.plot()
4 pyplot.show()
```

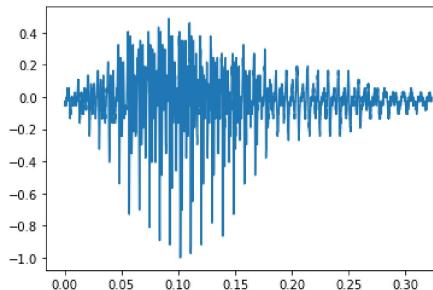


```
1 wave = thinkdsp.read_wave('/content/recording/1_nicolas_32.wav')
2 wave.play()
3 IPython.display.Audio('sound.wav')
```

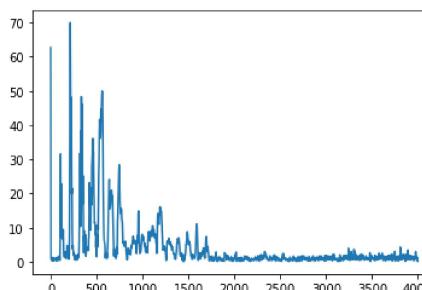
Writing sound.wav

0:00 / 0:00

```
1 segment = wave.segment(start=0)
2 segment.plot()
3 pyplot.show()
```



```
1 # Plot spectrum of wave
2 spectrum = wave.make_spectrum()
3 spectrum.plot()
4 pyplot.show()
```



## Creating dataframe

```
1 #Create a dataframe(name=df_audio) with two columns(path, label).
2 #You can get the label from the first letter of name.
3 #Eg: 0_jackson_0 --> 0
4 #0_jackson_43 --> 0
5

1 #info
2 df_audio.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
--- 
 0   path    2000 non-null   object  
 1   label   2000 non-null   object  
dtypes: object(2)
memory usage: 31.4+ KB
```

### Grader function 2

Saving... 

```
1 def grader_df():
2     flag_shape = df_audio.shape==(2000,2)
3     flag_columns = all(df_audio.columns==['path', 'label'])
4     list_values = list(df_audio.label.value_counts())
5     flag_label = len(list_values)==10
6     flag_label2 = all([i==200 for i in list_values])
7     final_flag = flag_shape and flag_columns and flag_label and flag_label2
8     return final_flag
9 grader_df()
```

True

```
1 df_audio = shuffle(df_audio, random_state=33)#don't change the random state
```

### Train and Validation split

```
1 #split the data into train and validation and save in X_train, X_test, y_train, y_test
2 #use stratify sampling
3 #use random state of 45
4 #use test size of 30%
5 X = df_audio.path
6 y = df_audio.label
7 X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,random_state=45,stratify=y)
```

### Grader function 3

```
1 def grader_split():
2     flag_len = (len(X_train)==1400) and (len(X_test)==600) and (len(y_train)==1400) and (len(y_test)==600)
3     values_ytrain = list(y_train.value_counts())
4     flag_ytrain = (len(values_ytrain)==10) and (all([i==140 for i in values_ytrain]))
5     values_ytest = list(y_test.value_counts())
6     flag_ytest = (len(values_ytest)==10) and (all([i==60 for i in values_ytest]))
7     final_flag = flag_len and flag_ytrain and flag_ytest
8     return final_flag
9 grader_split()
```

True

### Preprocessing

All files are in the "WAV" format. We will read those raw data files using the librosa

```
1 sample_rate = 22050
2 def load_wav(x, get_duration=True):
3     '''This return the array values of audio with sampling rate of 22050 and Duration'''
4     #loading the wav file with sampling rate of 22050
5     samples, sample_rate = librosa.load(x, sr=22050)
6     if get_duration:
7         duration = librosa.get_duration(samples, sample_rate)
8         return [samples, duration]
9     else:
10        return samples

1 #use load_wav function that was written above to get every wave.
2 #save it in X_train_processed and X_test_processed
3 # X_train_processed/X_test_processed should be dataframes with two columns(raw_data, duration) with same index of X_train/y_train
4 def get_full_path(x):
```

```
5     return "/content/recording"+x
6 #data[['c','d']] = [givetup(a) for a in data['b']]
7 X_train_processed = pd.DataFrame()
8 X_test_processed = pd.DataFrame()
9 X_train_processed[['raw_data', 'duration']] = [load_wav(get_full_path(x)) for x in X_train]
10 X_test_processed[['raw_data', 'duration']] = [load_wav(get_full_path(x)) for x in X_test]
```

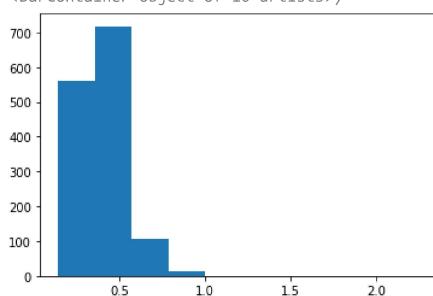
```
1 X_train_processed.head()
```

```
 raw_data duration
```

0	[-0.006193983, -0.0053070816, -0.0017616726, 0..., 0.58127
1	[-0.00013747877, -0.0002816826, -0.00045083734..., 0.24966
2	[-0.004207985, -0.0012800224, 0.002231387, 0.0..., 0.714785
3	[-0.000104876926, -9.907236e-05, -6.296669e-05..., 0.303401
4	[-0.0005594357, -0.0001103753, 0.00031156163, ..., 0.330522

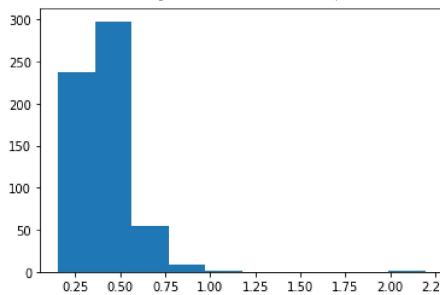
```
1 #plot the histogram of the duration for trian
2 plt.hist(X_train_processed.duration)
```

```
(array([561., 718., 107., 12., 0., 0., 0., 0., 1.]),
 array([0.1435374149659864, 0.3574603174603175, 0.5713832199546486,
 Saving...  
         1.992290249433107, 1.2131519274376419,
         1.40997732426304, 1.854920634920635,
         1.2766439909297], dtype=object),  
<BarContainer object of 10 artists>)
```



```
1 #plot the histogram of the duration for test
2 plt.hist(X_test_processed.duration)
```

```
(array([237., 298., 55., 8., 1., 0., 0., 0., 1.]),
 array([0.1564172335600907, 0.36036734693877553, 0.5643174603174603,
 0.7682675736961452, 0.97221768707483, 1.1761678004535148,
 1.3801179138321995, 1.5840680272108845, 1.7880181405895692,
 1.991968253968254, 2.195918367346939], dtype=object),
<BarContainer object of 10 artists>)
```



```
1 #print 0 to 100 percentile values with step size of 10 for train data duration.
2 for i in range(0,11,1):
3     print(f"{i*10}th percentile is {np.percentile(X_train_processed.duration.values, i*10)}")
```

```
0th percentile is 0.143537414965984
10th percentile is 0.25997732426303855
20th percentile is 0.298140589569161
30th percentile is 0.32884807256235826
40th percentile is 0.3573061224489796
50th percentile is 0.38770975056689344
60th percentile is 0.41365986394557824
70th percentile is 0.4454104308390022
80th percentile is 0.4823401360544218
90th percentile is 0.5528934240362812
100th percentile is 2.282766439909297
```

```
1 ##print 90 to 100 percentile values with step size of 1.
2 for i in range(0,11,1):
3     print(f"{i+90}th percentile is {np.percentile(X_train_processed.duration.values, i+90)}")
```

```
90th percentile is 0.5528934240362812
91th percentile is 0.5619047619047619
92th percentile is 0.5790748299319729
93th percentile is 0.5974403628117915
94th percentile is 0.610172335600907
95th percentile is 0.624079365079365
96th percentile is 0.6379374149659863
97th percentile is 0.6528367346938775
98th percentile is 0.6823900226757368
99th percentile is 0.7832580498866212
100th percentile is 2.282766439909297
```

## Grader function 4

```
1 def grader_processed():
2     flag_columns = (all(X_train_processed.columns==['raw_data', 'duration'])) and (all(X_test_processed.columns==['raw_data', 'duration']))
3     flag_shape = (X_train_processed.shape == (1400, 2)) and (X_test_processed.shape== (600,2))
4     return flag_columns and flag_shape
5 grader_processed()

True
```

Based on our analysis 99 percentile values are less than 0.8sec so we will limit maximum length of X\_train\_processed and X\_test\_processed to 0.8 sec. It is similar to pad\_sequence for a text dataset.

While loading the audio files, we are using sampling rate of 22050 so one sec will give array of length 22050. so, our maximum length is  $0.8 \times 22050 = 17640$  Pad with Zero if length of sequence is less than 17640 else Truncate the number.

Also create a masking vector for train and test.

masking vector value = 1 if it is real value, 0 if it is pad value. Masking vector data type must be bool.

```
1 max_length  = 17640

Saving... × (X_train_processed.raw_data, maxlen = max_length, padding = 'post',dtype = 'float32',truncating = 'post')
(X_test_processed.raw_data, maxlen = max_length,padding = 'post',dtype = 'float32',truncating = 'post')

3
4 X_train_mask = X_train_pad_seq != 0
5 X_test_mask = X_test_pad_seq != 0

1 ## as discussed above, Pad with Zero if length of sequence is less than 17640 else Truncate the number.
2 ## save in the X_train_pad_seq, X_test_pad_seq
3 ## also Create masking vector X_train_mask, X_test_mask
4
5 ## all the X_train_pad_seq, X_test_pad_seq, X_train_mask, X_test_mask will be numpy arrays mask vector dtype must be bool.
```

## Grader function 5

```
1 def grader_padoutput():
2     flag_padshape = (X_train_pad_seq.shape==(1400, 17640)) and (X_test_pad_seq.shape==(600, 17640)) and (y_train.shape==(1400,))
3     flag_maskshape = (X_train_mask.shape==(1400, 17640)) and (X_test_mask.shape==(600, 17640)) and (y_test.shape==(600,))
4     flag_dtype = (X_train_mask.dtype==bool) and (X_test_mask.dtype==bool)
5     return flag_padshape and flag_maskshape and flag_dtype
6 grader_padoutput()

True
```

### 1. Giving Raw data directly.

Now we have

Train data: X\_train\_pad\_seq, X\_train\_mask and y\_train

Test data: X\_test\_pad\_seq, X\_test\_mask and y\_test

We will create a LSTM model which takes this input.

Task:

1. Create an LSTM network which takes "X\_train\_pad\_seq" as input, "X\_train\_mask" as mask input. You can use any number of LSTM cells.

Please read LSTM documentation([https://www.tensorflow.org/api\\_docs/python/tf/keras/layers/LSTM](https://www.tensorflow.org/api_docs/python/tf/keras/layers/LSTM)) in tensorflow to know more about mask and also [https://www.tensorflow.org/guide/keras/masking\\_and\\_padding](https://www.tensorflow.org/guide/keras/masking_and_padding)

2. Get the final output of the LSTM and give it to Dense layer of any size and then give it to Dense layer of size 10(because we have 10 outputs) and then compile with the sparse categorical cross entropy( because we are not converting it to one hot vectors). Also check the datatype of class labels(y\_values) and make sure that you convert your class labels to integer datatype before fitting in the model.

3. While defining your model make sure that you pass both the input layer and mask input layer as input to lstm layer as follows

```
lstm_output = self.lstm(input_layer, mask=masking_input_layer)
```

4. Use tensorboard to plot the graphs of loss and metric(use custom micro F1 score as metric) and histograms of gradients. You can write your code for computing F1 score using this [link](#)

5. make sure that it won't overfit.

6. You are free to include any regularization

```
1 ## as discussed above, please write the architecture of the model.
2 ## you will have two input layers in your model (data input layer and mask input layer)
3 ## make sure that you have defined the data type of masking layer as bool
4
```

```
1 X_train_pad_seq = np.expand_dims(X_train_pad_seq,-1)
2 X_test_pad_seq = np.expand_dims(X_test_pad_seq,-1)
3 y_train = y_train.values.astype('int')
4 y_test = y_test.values.astype('int')
```

```

1 class F1ScoreCB(tf.keras.callbacks.Callback):
2     def __init__(self, train_data, test_data):
3         super().__init__()
4         self.train_data = train_data
5         self.test_data = test_data
6         self.history = {}
7         self.history['val_f1_score'] = []
8
9     def on_epoch_end(self, epochs, logs = {}):
10        train_preds = np.argmax(self.model.predict(self.train_data[0]), axis = -1)
11        train_f1_score = f1_score(self.train_data[1], train_preds, average='micro')
12        train_f1_score = np.round(train_f1_score, 4)
13
14        test_preds = np.argmax(self.model.predict(self.test_data[0]), axis = -1)
15        test_f1_score = f1_score(self.test_data[1], test_preds, average='micro')
16        test_f1_score = np.round(test_f1_score, 4)
17        self.history['val_f1_score'].append(test_f1_score)
18
19        print(f" - f1_score: {train_f1_score} - val_f1_score: {test_f1_score}")
20
21        writer1 = tf.summary.create_file_writer(logdir + '/train_f1_score')
22        writer2 = tf.summary.create_file_writer(logdir + '/validation_f1_score')
23
24        #writing to tensoboard
25        writer1.t()
26        Saving... × 'F1_Score', train_f1_score, step=epochs)
27
28
29        with writer2.as_default():
30            tf.summary.scalar('F1_Score', test_f1_score, step=epochs)
31        writer2.flush()

1 f1_score_cb = F1ScoreCB(([X_train_pad_seq, X_train_mask], y_train), ([X_test_pad_seq, X_test_mask], y_test))

1 #creating the first model
2 x= Input(shape = (17640,1)) # this will take inputs of 17640 time steps, each having dimension = 1
3 mask_indicator = Input(shape =(17640),dtype='bool') # to tell the model,ignore the pads
4 lstm_1 = LSTM(units=256)(x,mask=mask_indicator)
5 dense_1 = Dense(100,activation='relu')(lstm_1)
6 dense_2 = Dense(50,activation='relu')(dense_1)
7 dense_3 = Dense(20,activation='relu')(dense_2)
8 output = Dense(10, activation='softmax')(dense_3)
9 model_1 = Model(inputs=[x, mask_indicator], outputs=output)
10 opt = tf.keras.optimizers.Adam(learning_rate=0.00001)
11 model_1.compile(loss='sparse_categorical_crossentropy', optimizer=opt, metrics=['accuracy'])
12 model_1.summary()

Model: "model"

```

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 17640, 1)]	0	[]
input_2 (InputLayer)	[(None, 17640)]	0	[]
lstm (LSTM)	(None, 256)	264192	['input_1[0][0]', 'input_2[0][0]']
dense (Dense)	(None, 100)	25700	['lstm[0][0]']
dense_1 (Dense)	(None, 50)	5050	['dense[0][0]']
dense_2 (Dense)	(None, 20)	1020	['dense_1[0][0]']
dense_3 (Dense)	(None, 10)	210	['dense_2[0][0]']

```

Total params: 296,172
Trainable params: 296,172
Non-trainable params: 0

```

---

```

1 tensorboard = TensorBoard(log_dir = './logs')
2
3 checkpoint = ModelCheckpoint('best_model_1.h',save_best_only=True,mode='auto')

1 !rm -rf ./logs/
2 import datetime
3
4 logdir = os.path.join('logs', datetime.datetime.now().strftime('%Y%m%d-%H%M%S'))

1 history = model_1.fit([X_train_pad_seq, X_train_mask], y_train,
2                     validation_data = ([X_test_pad_seq, X_test_mask], y_test),
3                     batch_size = 16 , epochs = 10,callbacks=[f1_score_cb,checkpoint,tensorboard],verbose=1)

Epoch 1/10
44/44 [=====] - 11s 239ms/step
19/19 [=====] - 5s 243ms/step
- f1_score: 0.0964 - val_f1_score: 0.0867
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These func
88/88 [=====] - 88s 1s/step - loss: 2.3026 - accuracy: 0.0929 - val_loss: 2.3026 - val_accuracy: 0.0867
Epoch 2/10
44/44 [=====] - 11s 238ms/step
19/19 [=====] - 5s 257ms/step
- f1_score: 0.0993 - val_f1_score: 0.095

```

```

WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These func
88/88 [=====] - 86s 975ms/step - loss: 2.3026 - accuracy: 0.1014 - val_loss: 2.3026 - val_accuracy: 0.0950
Epoch 3/10
44/44 [=====] - 11s 239ms/step
19/19 [=====] - 5s 244ms/step
- f1_score: 0.1079 - val_f1_score: 0.1
88/88 [=====] - 75s 861ms/step - loss: 2.3026 - accuracy: 0.0950 - val_loss: 2.3026 - val_accuracy: 0.1000
Epoch 4/10
44/44 [=====] - 11s 238ms/step
19/19 [=====] - 5s 239ms/step
- f1_score: 0.0986 - val_f1_score: 0.0833
88/88 [=====] - 76s 864ms/step - loss: 2.3025 - accuracy: 0.0964 - val_loss: 2.3026 - val_accuracy: 0.0833
Epoch 5/10
44/44 [=====] - 11s 240ms/step
19/19 [=====] - 5s 241ms/step
- f1_score: 0.0964 - val_f1_score: 0.1033
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These func
88/88 [=====] - 75s 855ms/step - loss: 2.3025 - accuracy: 0.0929 - val_loss: 2.3026 - val_accuracy: 0.1033
Epoch 6/10
44/44 [=====] - 11s 242ms/step
19/19 [=====] - 5s 246ms/step
- f1_score: 0.0986 - val_f1_score: 0.105
88/88 [=====] - 64s 731ms/step - loss: 2.3025 - accuracy: 0.0900 - val_loss: 2.3026 - val_accuracy: 0.1050
Epoch 7/10
44/44 [=====] - 11s 239ms/step
19/19 [=====] - 5s 244ms/step
- f1_score: 0.1014 - val_f1_score: 0.0933
88/88 [=====] - 66s 747ms/step - loss: 2.3025 - accuracy: 0.0857 - val_loss: 2.3026 - val_accuracy: 0.0933
Saving...

```

=====] - 11s 238ms/step  
=====] - 5s 243ms/step

```

- f1_score: 0.1014 - val_f1_score: 0.095
88/88 [=====] - 66s 750ms/step - loss: 2.3025 - accuracy: 0.0986 - val_loss: 2.3026 - val_accuracy: 0.0950
Epoch 9/10
44/44 [=====] - 11s 239ms/step
19/19 [=====] - 5s 239ms/step
- f1_score: 0.1007 - val_f1_score: 0.0933
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These func
88/88 [=====] - 85s 971ms/step - loss: 2.3025 - accuracy: 0.0986 - val_loss: 2.3026 - val_accuracy: 0.0933
Epoch 10/10
44/44 [=====] - 11s 235ms/step
19/19 [=====] - 5s 247ms/step
- f1_score: 0.0921 - val_f1_score: 0.0967
88/88 [=====] - 76s 869ms/step - loss: 2.3025 - accuracy: 0.0936 - val_loss: 2.3026 - val_accuracy: 0.0967

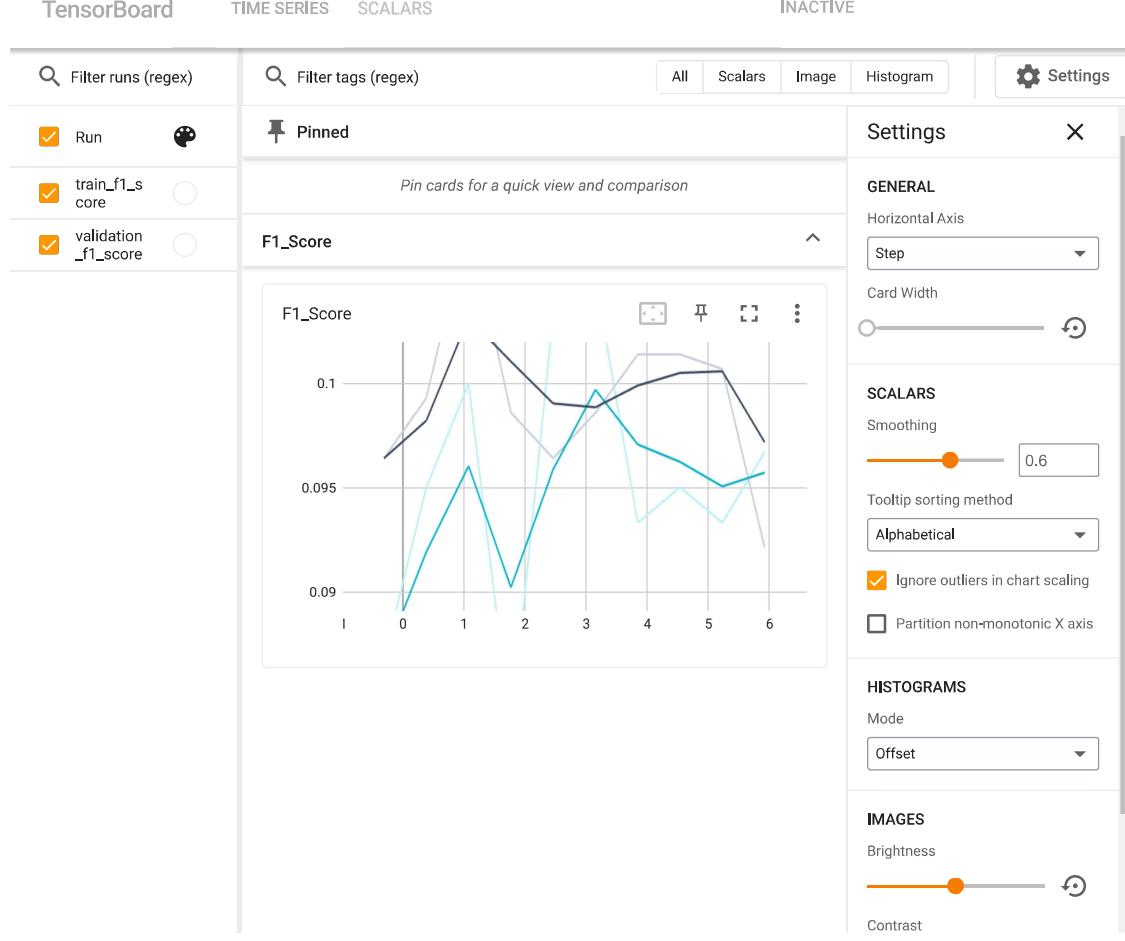
```

```

1 %reload_ext tensorboard
2 %tensorboard --logdir $logdir

```

Reusing TensorBoard on port 6006 (pid 7926), started 0:00:06 ago. (Use '!kill 7926' to kill it.)



We get a good fit model at epoch 5

2. Converting into spectrogram and giving spectrogram data as input

We can use librosa to convert raw data into spectrogram. A spectrogram shows the features in a two-dimensional representation with the intensity of a frequency at a point in time i.e we are converting Time domain to frequency domain. you can read more about this in <https://pnsn.org/spectrograms/what-is-a-spectrogram>

```
1 def convert_to_spectrogram(raw_data):
2     '''converting to spectrogram'''
3     spectrum = librosa.feature.melspectrogram(y=raw_data, sr=sample_rate, n_mels=64)
4     logmel_spectrum = librosa.power_to_db(S=spectrum, ref=np.max)
5     return logmel_spectrum

1 X_train_pad_seq = np.squeeze(X_train_pad_seq, axis = -1)
2 X_test_pad_seq = np.squeeze(X_test_pad_seq, axis = -1)

1 ##use convert_to_spectrogram and convert every raw sequence in X_train_pad_seq and X_test_pad_seq.
2 ## save those all in the X_train_spectrogram and X_test_spectrogram ( These two arrays must be numpy arrays)
3 X_train_spectrogram = np.array([convert_to_spectrogram(seq) for seq in X_train_pad_seq])
4 X_test_spectrogram = np.array([convert_to_spectrogram(seq) for seq in X_test_pad_seq])
```

## Grader function 6

```
1 def grader_spectrogram():
Saving... ✘ X program.shape==(1400,64, 35)) and (X_test_spectrogram.shape == (600, 64, 35))
4 grader_spectrogram()
True
```

Now we have

Train data: X\_train\_spectrogram and y\_train

Test data: X\_test\_spectrogram and y\_test

We will create a LSTM model which takes this input.

Task:

1. Create an LSTM network which takes "X\_train\_spectrogram" as input and has to return output at every time step.
2. Average the output of every time step and give this to the Dense layer of any size. (ex: Output from LSTM will be (None, time\_steps, features) average the output of every time step i.e, you should get (None,time\_steps) and then pass to dense layer )
3. give the above output to Dense layer of size 10( output layer) and train the network with sparse categorical cross entropy.
4. Use tensorboard to plot the graphs of loss and metric(use custom micro F1 score as metric) and histograms of gradients. You can write your code for computing F1 score using this [link](#)
5. make sure that it won't overfit.
6. You are free to include any regularization

```
1 # write the architecture of the model
2 #print model.summary and make sure that it is following point 2 mentioned above
3 x= Input(shape = (64,35,)) # this will take inputs of 64 time steps, each 35 dimensional
4 lstm_1 = LSTM(units=256,return_sequences=True)(x)
5 lstm_1_averaged = tf.math.reduce_mean(lstm_1, axis = -1)
6 dense_1 = Dense(256,activation='relu',kernel_initializer = 'he_normal')(lstm_1_averaged)
7 dense_2 = Dense(100,activation='relu')(dense_1)
8 dense_3 = Dense(50,activation='relu',kernel_initializer = 'he_normal')(dense_2)
9 output = Dense(10, activation='softmax')(dense_3)
10 model_2 = Model(inputs=x, outputs=output)
11 opt = tf.keras.optimizers.Adam() # usign the default learning rate
12 model_2.compile(loss='sparse_categorical_crossentropy', optimizer=opt, metrics=['accuracy'])
13 model_2.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
<hr/>		
input_1 (InputLayer)	[(None, 64, 35)]	0
lstm (LSTM)	(None, 64, 256)	299008
tf.math.reduce_mean (TFOpLambda)	(None, 64)	0
dense (Dense)	(None, 256)	16640
dense_1 (Dense)	(None, 100)	25700
dense_2 (Dense)	(None, 50)	5050
dense_3 (Dense)	(None, 10)	510
<hr/>		
Total params:	346,908	
Trainable params:	346,908	
Non-trainable params:	0	

```
1 #compile and fit your model.
2 #model2.fit([X_train_spectrogram],y_train_int,.....)
```

```
1 import datetime
2 !rm -rf ./logs/
```

```

3 logdir = os.path.join('logs', datetime.datetime.now().strftime('%Y%m%d-%H%M%S'))
4 tensorboard = TensorBoard(log_dir = './logs')
5 checkpoint = ModelCheckpoint('best_model_2.h', save_best_only=True, mode='auto')
6 f1_score_cb = F1ScoreCB((X_train_spectrogram, y_train), (X_test_spectrogram, y_test))

1 history = model_2.fit(X_train_spectrogram, y_train,
2                         validation_data = (X_test_spectrogram, y_test),
3                         batch_size = 16 , epochs = 30,verbose=1, callbacks=[f1_score_cb,checkpoint,tensorboard])

Epoch 1/30
6/88 [=>.....] - ETA: 0s - loss: 2.3025 - accuracy: 0.0625      WARNING:tensorflow:Callback method `on_train_batch_end` is slow compare
44/44 [=====] - 1s 3ms/step
19/19 [=====] - 0s 3ms/step
- f1_score: 0.3407 - val_f1_score: 0.2933
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These f
88/88 [=====] - 23s 98ms/step - loss: 2.1544 - accuracy: 0.1829 - val_loss: 1.8738 - val_accuracy: 0.2933
Epoch 2/30
44/44 [=====] - 0s 3ms/step
19/19 [=====] - 0s 3ms/step
- f1_score: 0.4693 - val_f1_score: 0.445
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These f
88/88 [=====] - 6s 72ms/step - loss: 1.6410 - accuracy: 0.4043 - val_loss: 1.4439 - val_accuracy: 0.4450
Epoch 3/30
44/44 [=====] - 0s 3ms/step
19/19 [=====] - 0s 3ms/step
- f1_score: 0.5879 - val_f1_score: 0.5933
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These f
Saving... X =====] - 7s 84ms/step - loss: 1.2649 - accuracy: 0.5464 - val_loss: 1.1156 - val_accuracy: 0.5933

44/44 [=====] - 0s 3ms/step
19/19 [=====] - 0s 3ms/step
- f1_score: 0.6743 - val_f1_score: 0.6717
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These f
88/88 [=====] - 7s 81ms/step - loss: 1.0782 - accuracy: 0.6036 - val_loss: 1.0012 - val_accuracy: 0.6717
Epoch 5/30
44/44 [=====] - 0s 3ms/step
19/19 [=====] - 0s 5ms/step
- f1_score: 0.7393 - val_f1_score: 0.7283
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These f
88/88 [=====] - 7s 80ms/step - loss: 0.8854 - accuracy: 0.6971 - val_loss: 0.8146 - val_accuracy: 0.7283
Epoch 6/30
44/44 [=====] - 0s 3ms/step
19/19 [=====] - 0s 3ms/step
- f1_score: 0.7293 - val_f1_score: 0.7183
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These f
88/88 [=====] - 6s 73ms/step - loss: 0.8287 - accuracy: 0.6936 - val_loss: 0.7695 - val_accuracy: 0.7183
Epoch 7/30
44/44 [=====] - 0s 3ms/step
19/19 [=====] - 0s 4ms/step
- f1_score: 0.7721 - val_f1_score: 0.7533
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These f
88/88 [=====] - 7s 82ms/step - loss: 0.7451 - accuracy: 0.7271 - val_loss: 0.6852 - val_accuracy: 0.7533
Epoch 8/30
44/44 [=====] - 0s 3ms/step
19/19 [=====] - 0s 3ms/step
- f1_score: 0.7993 - val_f1_score: 0.7467
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These f
88/88 [=====] - 7s 75ms/step - loss: 0.7105 - accuracy: 0.7350 - val_loss: 0.6593 - val_accuracy: 0.7467
Epoch 9/30
44/44 [=====] - 0s 5ms/step
19/19 [=====] - 0s 4ms/step
- f1_score: 0.8029 - val_f1_score: 0.78
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These f
88/88 [=====] - 8s 97ms/step - loss: 0.6537 - accuracy: 0.7657 - val_loss: 0.5610 - val_accuracy: 0.7800
Epoch 10/30
44/44 [=====] - 0s 3ms/step

```

```

1 %reload_ext tensorboard
2 %tensorboard --logdir $logdir

```

Show data download links

Ignore outliers in chart scaling

Tooltip sorting method: **default**

Filter tags (regular expressions supported)

F1\_Score

F1\_Score  
tag: F1\_Score

After converting the time domain speech recordings data to frequency domain, we are able to learn the classification in a much better manner.

We trained it for 30 epochs. Clearly it can be trained for an even larger epochs, we stopped due to computational limitation.

### 3. Data augmentation with raw features

Till now we have done with 2000 samples only. It is very less data. We are giving the process of generating augmented data below.

There are two types of augmentation:

1. time stretching - Time stretching either increases or decreases the length of the file. For time stretching we move the file 30% faster or slower

es the frequencies higher or lower. For pitch shifting we shift up or down one half-step.

Saving...

```
1 ## generating augmented data.
2 def generate_augmented_data(file_path):
3     augmented_data = []
4     samples = load_wav(file_path, get_duration=False)
5     for time_value in [0.7, 1, 1.3]:
6         for pitch_value in [-1, 0, 1]:
7             time_stretch_data = librosa.effects.time_stretch(samples, rate=time_value)
8             final_data = librosa.effects.pitch_shift(time_stretch_data, sr=sample_rate, n_steps=pitch_value)
9             augmented_data.append(final_data)
10    return augmented_data
```

```
1 temp_path = df_audio.iloc[0].path
2 aug_temp = generate_augmented_data('/content/recording/' + temp_path)
```

```
1 np.array(aug_temp).shape
```

```
(9,)
```

### Follow the steps

1. Split data 'df\_audio' into train and test (80-20 split)
2. We have 2000 data points(1600 train points, 400 test points)

```
1 X_train, X_test, y_train, y_test = train_test_split(df_audio['path'], df_audio['label'], random_state=45, test_size=0.2, stratify=df_audio['label'])
```

3. Do augmentation only on X\_train, pass each point of X\_train to generate\_augmented\_data function. After augmentation we will get 14400 train points. Make sure that you are augmenting the corresponding class labels (y\_train) also.
4. Preprocess your X\_test using load\_wav function.
5. Convert the augmented\_train\_data and test\_data to numpy arrays.
6. Perform padding and masking on augmented\_train\_data and test\_data.
7. After padding define the model similar to model 1 and fit the data

**Note** - While fitting your model on the augmented data for model 3 you might face Resource exhaust error. One simple hack to avoid that is save the augmented\_train\_data, augmented\_y\_train, test\_data and y\_test to Drive or into your local system. Then restart the runtime so that now you can train your model with full RAM capacity. Upload these files again in the new runtime session perform padding and masking and then fit your model.

```
1 X_train_augmented = []
2 y_train_augmented = []
3 for path, label in tqdm(zip(X_train.values, y_train.values), total=1600):
4     path = "/content/recording/" + path
5     augmented_datapoints = generate_augmented_data(path)
6     augmented_labels = [label] * 9
7     X_train_augmented.extend(augmented_datapoints)
8     y_train_augmented.extend(augmented_labels)
```

```
100%|██████████| 1600/1600 [05:41<00:00, 4.68it/s]
```

```

1 X_test_augmented = []
2 y_test_augmented = []
3 for path,label in tqdm(zip(X_test.values,y_test.values),total=400):
4     path = "/content/recording/" + path
5     augmented_datapoints = generate_augmented_data(path)
6     augmented_labels = [label]*9
7     X_test_augmented.extend(augmented_datapoints)
8     y_test_augmented.extend(augmented_labels)

100%|██████████| 400/400 [01:32<00:00,  4.32it/s]

1 X_train_augmented = np.array(X_train_augmented)
2 X_test_augmented = np.array(X_test_augmented)

1 y_train_augmented = np.array(y_train_augmented,dtype='int')
2 y_test_augmented = np.array(y_test_augmented,dtype='int')

1 x= Input(shape = (17640,1)) # this will take inputs of 17640 time steps, each having dimension = 1
2 mask_indicator = Input(shape =(17640),dtype='bool') # to tell the model, ignore the pads
3 lstm_1 = LSTM(units=256)(x,mask=mask_indicator)
4 dense_1 = Dense(100,activation='relu')(lstm_1)
5 dense_2 = Dense(50,activation='relu')(dense_1)
6 dense_3 = Dense(20,activation='relu')(dense_2)
Saving...                               softmax')(dense_3)
9 opt = tf.keras.optimizers.Adam(learning_rate=0.00001)
10 model_3.compile(loss='sparse_categorical_crossentropy', optimizer=opt, metrics=['accuracy'])
11 model_3.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
<hr/>			
input_1 (InputLayer)	[(None, 17640, 1)]	0	[]
input_2 (InputLayer)	[(None, 17640)]	0	[]
lstm (LSTM)	(None, 256)	264192	['input_1[0][0]', 'input_2[0][0]']
dense (Dense)	(None, 100)	25700	['lstm[0][0]']
dense_1 (Dense)	(None, 50)	5050	['dense[0][0]']
dense_2 (Dense)	(None, 20)	1020	['dense_1[0][0]']
dense_3 (Dense)	(None, 10)	210	['dense_2[0][0]']
<hr/>			
Total params: 296,172			
Trainable params: 296,172			
Non-trainable params: 0			

---

```

1 !rm -rf ./logs/
2 import datetime
3 logdir = os.path.join('logs', datetime.datetime.now().strftime('%Y%m%d-%H%M%S'))

1 import pickle
2 X_train_augmented = open("/content/drive/MyDrive/Copy of X_train_augmented (1).pkl", "rb")
3 X_test_augmented = open("/content/drive/MyDrive/Copy of X_test_augmented.pkl","rb")
4 y_train_augmented = open("/content/drive/MyDrive/Copy of y_train_augmented.pkl","rb")
5 y_test_augmented = open("/content/drive/MyDrive/Copy of y_test_augmented.pkl","rb")
6 X_train_augmented = pickle.load(X_train_augmented)
7 X_test_augmented = pickle.load(X_test_augmented)
8 y_train_augmented = pickle.load(y_train_augmented)
9 y_test_augmented = pickle.load(y_test_augmented)

1 max_length = 17640
2 X_train_augmented_padded = pad_sequences(X_train_augmented, maxlen = max_length,dtype = 'float32', padding = 'post',truncating = 'post')
3 X_test_augmented_padded = pad_sequences(X_test_augmented, maxlen = max_length, dtype = 'float32', padding = 'post', truncating = 'post')
4
5 X_train_augmented_mask = X_train_augmented_padded != 0
6 X_test_augmented_mask = X_test_augmented_padded != 0

1 X_train_augmented_padded = np.expand_dims(X_train_augmented_padded, -1)
2 X_test_augmented_padded = np.expand_dims(X_test_augmented_padded, -1)

1 tensorboard = TensorBoard(log_dir = './logs')
2 checkpoint = ModelCheckpoint('best_model_3.h',save_best_only=True,mode='auto')
3 f1_score_cb = F1ScoreCB(([X_train_augmented_padded, X_train_augmented_mask], y_train_augmented), ([X_test_augmented_padded, X_test_augmented_mask], y_test_augmented))

1 history = model_3.fit([X_train_augmented_padded, X_train_augmented_mask], y_train_augmented,
2                     validation_data = ([X_test_augmented_padded, X_test_augmented_mask], y_test_augmented),
3                     batch_size = 32 , epochs = 2 , callbacks=[f1_score_cb,checkpoint,tensorboard],verbose=1)

Epoch 1/2
450/450 [=====] - 125s 274ms/step
113/113 [=====] - 33s 275ms/step
- f1_score: 0.1 - val_f1_score: 0.1
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These func
450/450 [=====] - 532s 1s/step - loss: 2.3026 - accuracy: 0.0973 - val_loss: 2.3026 - val_accuracy: 0.1000

```

```

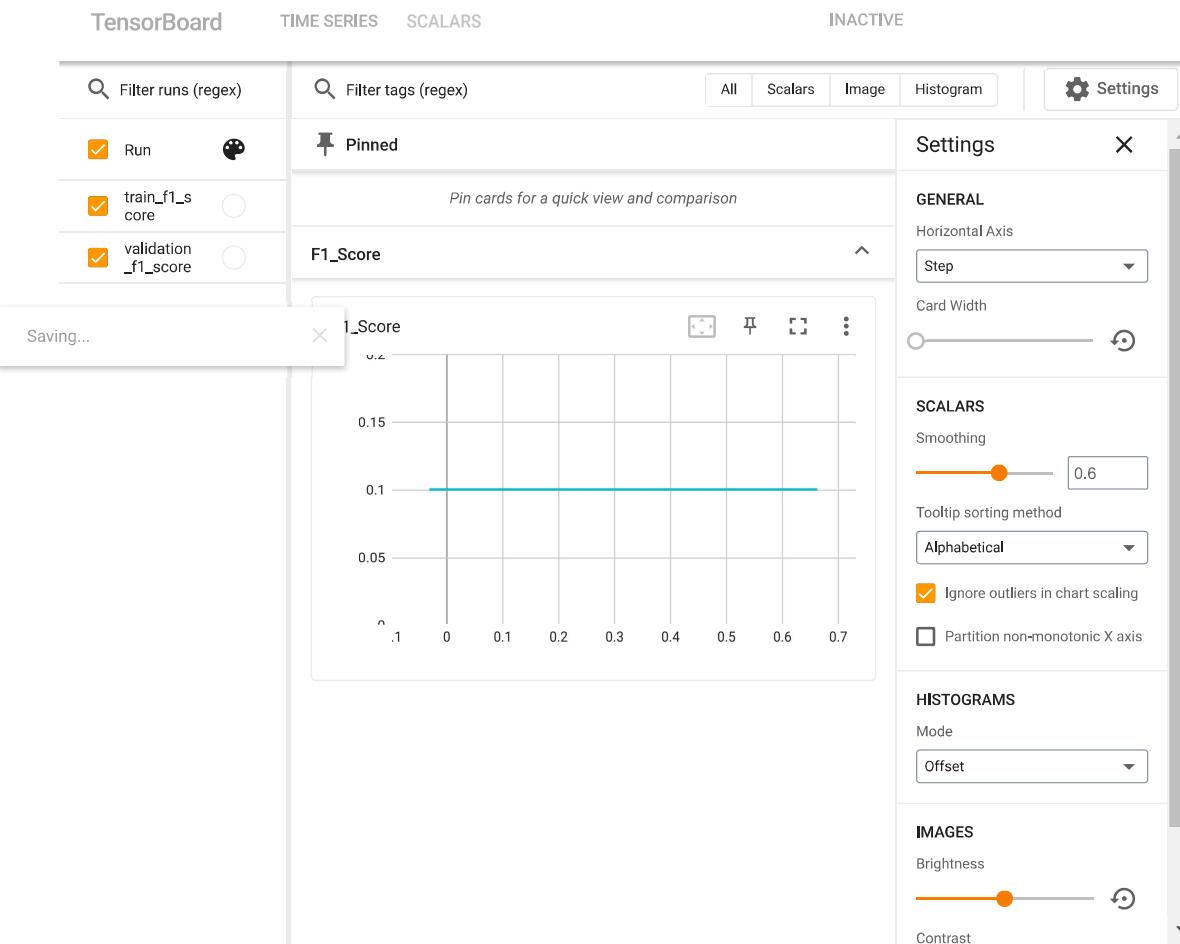
Epoch 2/2
450/450 [=====] - 123s 274ms/step
113/113 [=====] - 31s 273ms/step
- f1_score: 0.1 - val_f1_score: 0.1
WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These func
450/450 [=====] - 504s 1s/step - loss: 2.3026 - accuracy: 0.0970 - val_loss: 2.3026 - val_accuracy: 0.1000

```

```

1 %reload_ext tensorboard
2 %tensorboard --logdir $logdir

```



#### 4. Data augmentation with spectrogram data

1. use convert\_to\_spectrogram and convert the padded data from train and test data to spectrogram data.
2. The shape of train data will be 14400 x 64 x 35 and shape of test\_data will be 400 x 64 x35
3. Define the model similar to model 2 and fit the data

```

1 X_train_augmented_padded = np.squeeze(X_train_augmented_padded, axis = -1)
2 X_test_augmented_padded = np.squeeze(X_test_augmented_padded, axis = -1)

1 X_train_spectrogram_augmented = np.array([convert_to_spectrogram(seq) for seq in X_train_augmented_padded])
2 X_test_spectrogram_augmented = np.array([convert_to_spectrogram(seq) for seq in X_test_augmented_padded])

1 x= Input(shape = (64,35,)) # this will take inputs of 64 time steps, each 35 dimensional
2 lstm_1 = LSTM(units=256,return_sequences=True)(x)
3 lstm_1_averaged = tf.math.reduce_mean(lstm_1, axis = -1)
4 dense_1 = Dense(256,activation='relu',kernel_initializer = 'he_normal')(lstm_1_averaged)
5 dense_2 = Dense(100,activation='relu')(dense_1)
6 dense_3 = Dense(50,activation='relu',kernel_initializer = 'he_normal')(dense_2)
7 output = Dense(10, activation='softmax')(dense_3)
8 model_4 = Model(inputs=x, outputs=output)
9 opt = tf.keras.optimizers.Adam(learning_rate=0.00005)
10 model_4.compile(loss='sparse_categorical_crossentropy', optimizer=opt, metrics=['accuracy'])
11 model_4.summary()

```

Model: "model\_1"

Layer (type)	Output Shape	Param #
<hr/>		
input_3 (InputLayer)	[(None, 64, 35)]	0
lstm_1 (LSTM)	(None, 64, 256)	299008
tf.math.reduce_mean (TFOpLa mbda)	(None, 64)	0
dense_4 (Dense)	(None, 256)	16640
dense_5 (Dense)	(None, 100)	25700

```
dense_6 (Dense)          (None, 50)      5050
```

```
dense_7 (Dense)          (None, 10)       510
```

```
=====
Total params: 346,908
Trainable params: 346,908
Non-trainable params: 0
```

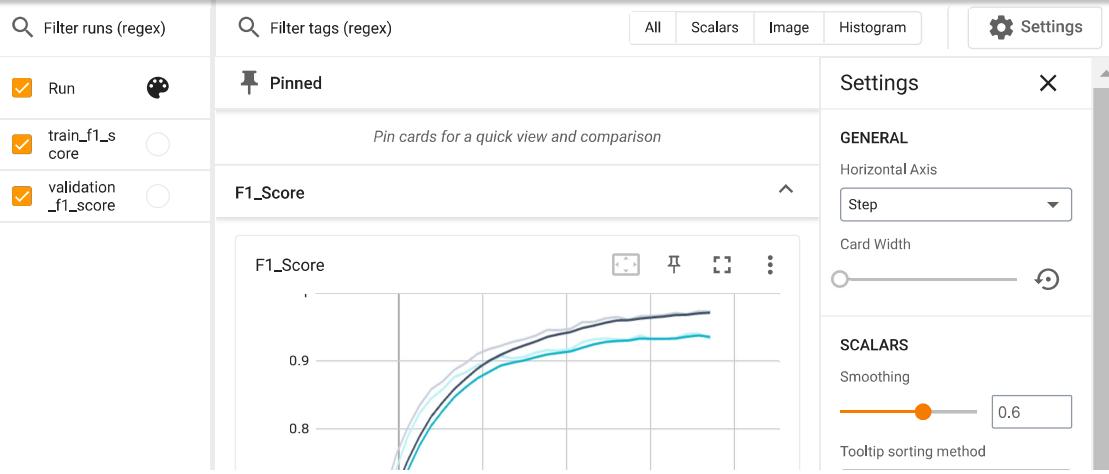
```
1 !rm -rf ./logs/
2 logdir = os.path.join('logs', datetime.datetime.now().strftime('%Y%m%d-%H%M%S'))
3 tensorboard = TensorBoard(log_dir = './logs')
4 checkpoint = ModelCheckpoint('best_model_4.h', save_best_only=True, mode='auto')
5 f1_score_cb = F1ScoreCB((X_train_spectrogram_augmented, y_train_augmented), (X_test_spectrogram_augmented, y_test_augmented))

1 history = model_4.fit(X_train_spectrogram_augmented, y_train_augmented,
2                         validation_data = (X_test_spectrogram_augmented, y_test_augmented),
3                         batch_size = 16 , epochs = 30, verbose=1, callbacks=[f1_score_cb,checkpoint,tensorboard])

Epoch 1/30
1/900 [........................] - ETA: 2:43:06 - loss: 2.3000 - accuracy: 0.0625WARNING:tensorflow:Callback method `on_train_batch_end` is slow compare
450/450 [=====] - 3s 4ms/step
113/113 [=====] - 1s 3ms/step
- f1_score: 0.5844
Saving... ✘ Actions such as lstm_cell_1_layer_call_fn, lstm_cell_1_layer_call_and_return_conditional_losses while saving (showing 2 of 2). The
900/900 [=====] - 29s 20ms/step - loss: 1.8313 - accuracy: 0.3875 - val_loss: 1.3196 - val_accuracy: 0.5844

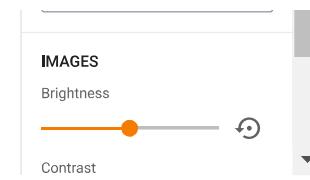
Epoch 2/30
450/450 [=====] - 1s 3ms/step
113/113 [=====] - 0s 4ms/step
- f1_score: 0.7015 - val_f1_score: 0.6883
WARNING:absl:Found untraced functions such as lstm_cell_1_layer_call_fn, lstm_cell_1_layer_call_and_return_conditional_losses while saving (showing 2 of 2). The
900/900 [=====] - 17s 19ms/step - loss: 1.0905 - accuracy: 0.6500 - val_loss: 0.9327 - val_accuracy: 0.6883
Epoch 3/30
450/450 [=====] - 2s 4ms/step
113/113 [=====] - 1s 5ms/step
- f1_score: 0.7609 - val_f1_score: 0.7394
WARNING:absl:Found untraced functions such as lstm_cell_1_layer_call_fn, lstm_cell_1_layer_call_and_return_conditional_losses while saving (showing 2 of 2). The
900/900 [=====] - 19s 21ms/step - loss: 0.8248 - accuracy: 0.7261 - val_loss: 0.7493 - val_accuracy: 0.7394
Epoch 4/30
450/450 [=====] - 1s 3ms/step
113/113 [=====] - 1s 5ms/step
- f1_score: 0.8014 - val_f1_score: 0.7881
WARNING:absl:Found untraced functions such as lstm_cell_1_layer_call_fn, lstm_cell_1_layer_call_and_return_conditional_losses while saving (showing 2 of 2). The
900/900 [=====] - 16s 18ms/step - loss: 0.6644 - accuracy: 0.7781 - val_loss: 0.6108 - val_accuracy: 0.7881
Epoch 5/30
450/450 [=====] - 2s 3ms/step
113/113 [=====] - 0s 3ms/step
- f1_score: 0.8333 - val_f1_score: 0.8219
WARNING:absl:Found untraced functions such as lstm_cell_1_layer_call_fn, lstm_cell_1_layer_call_and_return_conditional_losses while saving (showing 2 of 2). The
900/900 [=====] - 17s 19ms/step - loss: 0.5545 - accuracy: 0.8126 - val_loss: 0.5278 - val_accuracy: 0.8219
Epoch 6/30
450/450 [=====] - 1s 3ms/step
113/113 [=====] - 0s 3ms/step
- f1_score: 0.8574 - val_f1_score: 0.8439
WARNING:absl:Found untraced functions such as lstm_cell_1_layer_call_fn, lstm_cell_1_layer_call_and_return_conditional_losses while saving (showing 2 of 2). The
900/900 [=====] - 17s 19ms/step - loss: 0.4769 - accuracy: 0.8424 - val_loss: 0.4616 - val_accuracy: 0.8439
Epoch 7/30
450/450 [=====] - 1s 3ms/step
113/113 [=====] - 0s 3ms/step
- f1_score: 0.8693 - val_f1_score: 0.8575
WARNING:absl:Found untraced functions such as lstm_cell_1_layer_call_fn, lstm_cell_1_layer_call_and_return_conditional_losses while saving (showing 2 of 2). The
900/900 [=====] - 16s 18ms/step - loss: 0.4195 - accuracy: 0.8613 - val_loss: 0.4120 - val_accuracy: 0.8575
Epoch 8/30
450/450 [=====] - 1s 3ms/step
113/113 [=====] - 0s 4ms/step
- f1_score: 0.8862 - val_f1_score: 0.8756
WARNING:absl:Found untraced functions such as lstm_cell_1_layer_call_fn, lstm_cell_1_layer_call_and_return_conditional_losses while saving (showing 2 of 2). The
900/900 [=====] - 16s 18ms/step - loss: 0.3748 - accuracy: 0.8767 - val_loss: 0.3696 - val_accuracy: 0.8756
Epoch 9/30
450/450 [=====] - 2s 5ms/step
113/113 [=====] - 1s 4ms/step
- f1_score: 0.8965 - val_f1_score: 0.8828
WARNING:absl:Found untraced functions such as lstm_cell_1_layer_call_fn, lstm_cell_1_layer_call_and_return_conditional_losses while saving (showing 2 of 2). The
900/900 [=====] - 16s 17ms/step - loss: 0.3367 - accuracy: 0.8876 - val_loss: 0.3492 - val_accuracy: 0.8828
Epoch 10/30
450/450 [=====] - 1s 3ms/step
```

```
1 %reload_ext tensorboard
2 %tensorboard --logdir $logdir
```



Saving...

1. Using the spectrogram, makes our life easier for classifying spoken digit recognition.
2. The reason being, different digits when spoken have different frequency/pitch range.
3. Also, the time domain data is more prone to low-frequency noise, which gets shifted to high frequency region in frequency domain.
4. This is evident in model\_2 and model\_4 f1\_score vs epoch plot.



✓ 4s completed at 8:23 AM